

CS 231A CA Session: Problem Set 4 Review

Kevin Chen
May 13, 2016

PS4 Outline

- Problem 1: Viewpoint estimation
- Problem 2: Segmentation
 - Meanshift segmentation
 - Normalized cut

Problem 1: Viewpoint Estimation

Viewpoint estimation using Bag of Words

- Goal: Perform image classification
- Steps:
 - Build dictionary of codewords
 - Build features based on codewords
 - Train an SVM classifier on the features and labels

Viewpoint estimation using Bag of Words

- Goal: Perform image classification
- Steps:
 - **Build dictionary of codewords**
 - Build features based on codewords
 - Train an SVM classifier on the features and labels

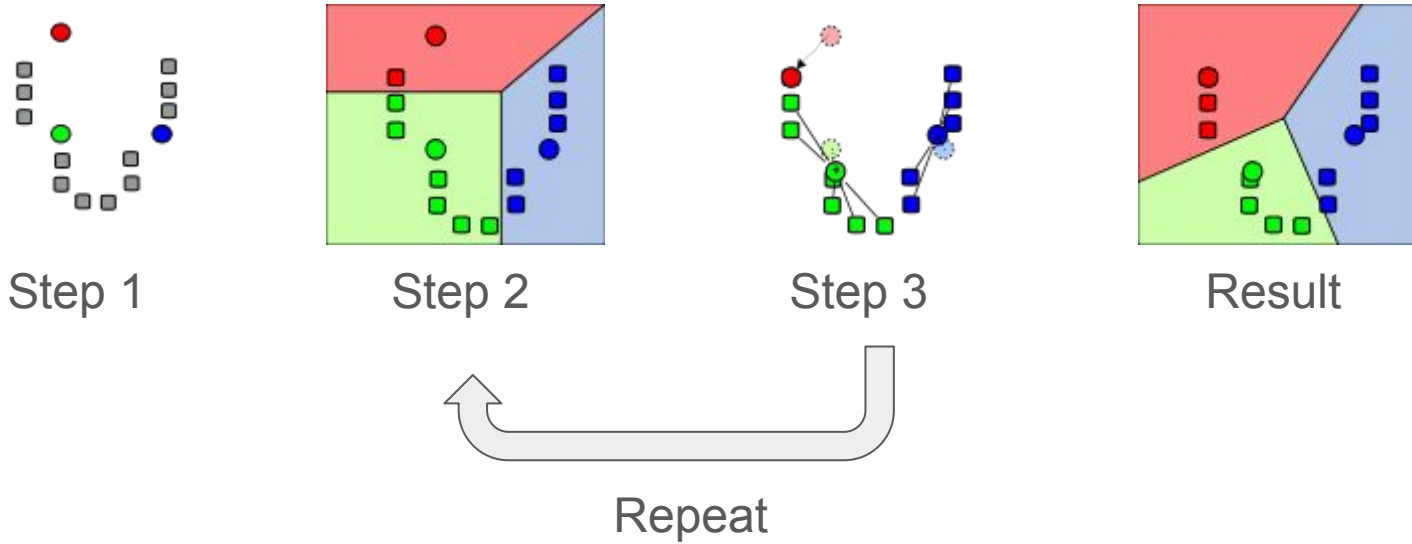
Build dictionary of codewords

- Goal: compute “words” for our bag of words model
 - Will be useful for computing a feature vector for an image
- Basic idea: perform k-means on dense SIFT features
- Cluster centers form the “words” in the bag of words model
- Algorithm
 - Subsample images in train set
 - Compute dense SIFT features
 - Find cluster centers by running k-means

K-means clustering

- Fixed number of clusters (k)
- Algorithm
 - Given an initial set of k means
 - Repeat until convergence
 - Assign each point to a i th cluster if it is closest to the i th mean (euclidean distance)
 - Recompute means
- Use VLFeat

k-means clustering



Viewpoint estimation using Bag of Words

- Goal: Perform image classification
- Steps:
 - Build dictionary of codewords
 - **Build features based on codewords**
 - Train an SVM classifier on the features and labels

Build features based on codewords

- Goal: Extract a feature from each image
- Algorithm
 - Extract dense SIFT features for every image in train set
 - For each image, build pyramid of dxd sub-images
 - For each pyramid depth
 - For every dense SIFT feature, find nearest codeword according to EUCLIDEAN distance
 - Build histogram of codewords
 - Concatenate sub-image histograms to build the final feature vector for the image

Viewpoint estimation using Bag of Words

- Goal: Perform image classification
- Steps:
 - Build dictionary of codewords
 - Build features based on codewords
 - **Train an SVM classifier on the features and labels**

Problem 2: Segmentation

Segmentation



Meanshift

Meanshift segmentation

- Compute features for each pixel, e.g. (u, v, R, G, B)
- Keep track of visited pixels (initialize all to unvisited)
- Keep track of cluster centers (initially there are none)
- Repeat until all pixels have been visited
 - Pick a random, unvisited pixel and get its corresponding feature vector
 - Perform meanshift procedure to get cluster center
 - Make sure to update visited pixels
 - Update list of cluster centers
 - Check if the newly computed cluster center is sufficiently close to any previous cluster centers
 - If so, ignore newly computed cluster center (say it's part of a previous cluster)
 - If not, add new cluster to cluster list
- Assign pixels to nearest cluster (in feature space)

Meanshift procedure

- Repeat until convergence
 - Find all features within bandwidth distance from the current feature vector
 - Compute mean of these features
 - All features with bandwidth circle (at any iteration) should have their corresponding pixels marked as visited
 - Update “current feature vector” as the newly computed mean
- Return the final mean as a new cluster center

Normalized Cut

Motivation

- Graph can be constructed with nodes as pixels and edges between pixels
 - Weights are associated with each edge
- Typically want to partition graph into two parts
 - Find partition corresponding to minimum cut cost
- Often ends up choosing cuts that partition small sets of nodes
- Use norm cut instead
 - Norm cut - cost is computed as a fraction of weights of edges connecting to the partition
- See paper for details

Basic ideas

- Vector \mathbf{x} defines a partition
 - $\mathbf{x}_i = 1$ if node i is in subset A
 - $\mathbf{x}_i = -1$ if node i is in subset B
- We want to solve the following problem:

$$\min_x \text{ncut}(x, W, D) :$$

- Solve for partition by trying to solve an equivalent problem

$$\min_x \text{ncut}(x, W, D) = \min_y \frac{y^T(W - D)y}{y^T D y}$$

$$\text{subject to } y_i \in \{1, -1\}, y^T D \mathbf{1} = 0$$

- \mathbf{y} is a reformulation of the partition indicator vector \mathbf{x}
 - We can retrieve A, B from \mathbf{y}

We ask you to implement...

- The computation of norm cut cost
- A recursive algorithm for partitioning the graph

We ask you to implement...

- **The computation of norm cut cost**
- A recursive algorithm for partitioning the graph

Computation of norm cut

- Given a partition value ϵ and vector \mathbf{v} , we can compute \mathbf{x}
 - If $\mathbf{v}_i > \epsilon$, then $\mathbf{x}_i = 1$
 - If $\mathbf{v}_i \leq \epsilon$, then $\mathbf{x}_i = -1$
- Use definitions to compute \mathbf{D} , \mathbf{k} , \mathbf{b} , and \mathbf{y}
- Cost can then be computed by

$$\text{ncut} = \frac{\mathbf{y}^T (\mathbf{D} - \mathbf{W}) \mathbf{y}}{\mathbf{y}^T \mathbf{D} \mathbf{y}}$$

Definitions

- Given a weight matrix \mathbf{W} in which $\mathbf{W}(i, j)$ corresponds to edge weight between node i and node j
- Compute diagonal matrix \mathbf{D} in which $\mathbf{D}(i, i)$ corresponds to sum of weights connecting to node i
- Vector \mathbf{x} defines a partition
 - $\mathbf{x}_i = 1$ if node i is in subset A
 - $\mathbf{x}_i = -1$ if node i is in subset B
- Others
 - Define \mathbf{k} to be sum of edge weights which connect to nodes in A
 - Define \mathbf{b} to be ratio of edge weights in A to edge weights in B

We ask you to implement...

- The computation of norm cut cost
- **A recursive algorithm for partitioning the graph**

Recursive algorithm for partitioning the graph

- Compute **D**
- Approximately solve for **y** in the problem shown before

$$\min_x \text{ncut}(x, W, D) = \min_y \frac{y^T(W - D)y}{y^T D y}$$

$$\text{subject to } y_i \in \{1, -1\}, y^T D \mathbf{1} = 0$$

- To do this, solve generalized eigensystem
 $(D - W)y = \lambda D y$

by finding the eigenvector **v** associated with the second smallest eigenvalue of the generalized eigensystem

- This does not satisfy first constraint, so find best partition point using *fminsearch* on norm cut cost
- Once a partition point is found, we can get A and B
- Recurse on partitions A and B unless end recursion conditions are true

When ending recursion...

- Suppose we want to partition V into A and B
- End recursion conditions
 - Norm cust cost is above threshold
 - Proposed A partition is too small
 - Proposed B partition is too small
- If any of the above are true...
 - Return V as the segmentation

Questions?